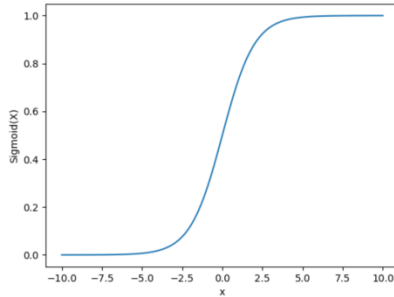



Experiment No. 3																	
BE (AI&DS)	ROLL NO : 9742																
Date of Implementation: 06/08/2024																	
Aim: To observe the impact of different activation functions on the outcome of the neural network.																	
Programming Language Used : Python																	
Upon completion of this experiment, students will be able to																	
L01 : Implement basic neural network model for a given a problem.																	
<table border="1" style="width: 100%; border-collapse: collapse;"> <tr> <td style="width: 33%; padding: 5px;">Indicator</td> <td style="width: 33%;"></td> <td style="width: 33%;"></td> </tr> <tr> <td style="padding: 5px;">Timeline Maintains submission deadline (1)</td> <td style="padding: 5px;">On time (1)</td> <td style="padding: 5px;">Otherwise (0)</td> </tr> <tr> <td style="padding: 5px;">Completion and Organization (2)</td> <td style="padding: 5px;">Completed in LAB (2 )</td> <td style="padding: 5px;">Otherwise(1)</td> </tr> <tr> <td style="padding: 5px;">Analysis of output and conclusion(2)</td> <td style="padding: 5px;">Properly done (2)</td> <td style="padding: 5px;">Otherwise (0)</td> </tr> <tr> <td style="padding: 5px;">Viva (10)</td> <td></td> <td></td> </tr> </table>			Indicator			Timeline Maintains submission deadline (1)	On time (1)	Otherwise (0)	Completion and Organization (2)	Completed in LAB (2 )	Otherwise(1)	Analysis of output and conclusion(2)	Properly done (2)	Otherwise (0)	Viva (10)		
Indicator																	
Timeline Maintains submission deadline (1)	On time (1)	Otherwise (0)															
Completion and Organization (2)	Completed in LAB (2 )	Otherwise(1)															
Analysis of output and conclusion(2)	Properly done (2)	Otherwise (0)															
Viva (10)																	
<p><b>Assessment Marks :</b></p> <table border="1" style="width: 100%; border-collapse: collapse;"> <tr> <td style="width: 80%; padding: 5px;">Timeline(1)</td> <td style="width: 20%;"></td> </tr> <tr> <td style="padding: 5px;">Completion and Organization (2)</td> <td></td> </tr> <tr> <td style="padding: 5px;">Analysis of output and conclusion(2)</td> <td></td> </tr> <tr> <td style="padding: 5px;">Viva (10)</td> <td></td> </tr> <tr> <td style="padding: 5px;">Total (15)</td> <td></td> </tr> </table>			Timeline(1)		Completion and Organization (2)		Analysis of output and conclusion(2)		Viva (10)		Total (15)						
Timeline(1)																	
Completion and Organization (2)																	
Analysis of output and conclusion(2)																	
Viva (10)																	
Total (15)																	

<b>EXPERIMENT</b>	<b>3</b>
<b>Aim</b>	To observe the impact of different activation functions on the outcome of the neural network.
<b>Tools</b>	PYTHON
<b>Theory</b>	<p>An activation function in the context of neural networks is a mathematical function applied to the output of a neuron. The purpose of an activation function is to introduce non-linearity into the model, allowing the network to learn and represent complex patterns in the data. The activation function decides whether a neuron should be activated or not by calculating the weighted sum and further adding bias to it.</p> <p><b>Variants of Activation Function</b></p> <p><b>Linear Function</b>  Linear function has the equation similar to as of a straight line i.e. <math>y = x</math>. No matter how many layers we have, if all are linear in nature, the final activation function of last layer is nothing but just a linear function of the input of first layer.</p> <p><b>Sigmoid Function</b></p>  <p>It is a function which is plotted as 'S' shaped graph. Equation : <math>A = 1/(1 + e^{-x})</math> Nature : Non-linear. Value Range : 0 to 1. Usually used in output layer of a binary classification, where result is either 0 or 1, as value for sigmoid function lies between 0 and 1 only so, result can be predicted easily to be 1 if value is greater than 0.5 and 0 otherwise.</p>

## Tanh Function



∂G

The activation that works almost always better than sigmoid function is Tanh function also known as Tangent Hyperbolic function. It's actually mathematically shifted version of the sigmoid function.

$$f(x) = \tanh(x) = 2 / (1 + e^{-2x}) - 1$$

OR

$$\tanh(x) = 2 * \text{sigmoid}(2x) - 1$$


Value Range :- -1 to +1 Nature :- non-linear

## RELU Function

It Stands for *Rectified linear unit*. It is the most widely used activation function. Chiefly implemented in *hidden layers* of Neural network.

Equation :-  $A(x) = \max(0, x)$ . It gives an output x if x is positive and 0 otherwise. Value Range :-  $[0, \infty)$  Nature :- non-linear ReLu is less computationally expensive than tanh and sigmoid because it involves simpler mathematical operations.

## Softmax Function



The softmax function is also a type of sigmoid function but is handy when we are trying to handle multi- class classification problems. Nature :- non-linear. The softmax function was commonly found in the output layer of image classification problems.

Implementati  
on

- 1: Import the necessary libraries. (tensorflow, keras, Sequential, Flatten, Dense, Activation)
- 2: Download the dataset.  
CIFAR-10 dataset (For few students)

	<p>Fashion MNIST dataset (For few students)</p> <p>3: Understand the structure of the dataset  4: Visualize the data.  5: Form the Input, hidden, and output layers. The Sequential model allows us to create models layer-by-layer as we need in a multi-layer perceptron and is limited to single-input, single-output stacks of layers. Flatten flattens the input provided without affecting the batch size. For example, If inputs are shaped (batch_size,) without a feature axis, then flattening adds an extra channel dimension and output shape is (batch_size, 1). Activation is for using the different activation function. The first two Dense layers are used to make a fully connected model and are the hidden layers. The last Dense layer is the output layer which contains neurons that decide which category the image belongs to.  6: Compile the model.  7: Fit the model.  8: Find Accuracy of the model.</p> <p>try to use different activation functions like linear, sigmoid, tanh, RELU, leaky RELU and softmax and note the impact on the accuracy model.</p>
Conclusion	<p>Discuss the impact of activation function on accuracy.</p> <p>In this experiment, we evaluated the performance of various activation functions in a neural network using the Fashion MNIST dataset. The aim was to understand how different activation functions impact the accuracy of the model.</p> <p>Results:</p> <ul style="list-style-type: none"> <li>• <b>Linear Activation:</b> Achieved an accuracy of 82.55%. The linear activation function provided a lower accuracy compared to the non-linear functions, demonstrating its limitations in capturing complex patterns due to its lack of non-linearity.</li> <li>• <b>Sigmoid Activation:</b> Obtained an accuracy of 87.75%. The sigmoid function, known for its non-linear characteristics, performed well in capturing non-linear relationships, resulting in the highest accuracy among the tested functions.</li> <li>• <b>Tanh Activation:</b> Recorded an accuracy of 87.5%. The tanh function also performed well, slightly below sigmoid, but still effectively handled non-linearity and provided strong performance.</li> </ul>

	<ul style="list-style-type: none"><li>• <b>ReLU Activation:</b> Achieved an accuracy of 87.14%. The ReLU function, which is computationally efficient and introduces non-linearity, performed similarly to tanh and sigmoid, though slightly lower in accuracy.</li><li>• <b>Softmax Activation:</b> Resulted in an accuracy of 77.73%. The softmax function, used in the output layer for multi-class classification, was less effective in this case when used in the hidden layers, reflecting its suitability primarily for output layers rather than hidden layers.</li></ul>
--	--

## Implementation:

```
[1]: import tensorflow as tf
      from tensorflow import keras
      from keras.models import Sequential
      from keras.layers import Flatten, Dense, Activation
      from tensorflow.keras.datasets import fashion_mnist
```

```
[2]: # Step 2: Download and prepare the Fashion MNIST dataset
      (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-labels-idx1-ubyte.gz
29515/29515          0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-images-idx3-ubyte.gz
26421880/26421880    0s
0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-labels-idx1-ubyte.gz
5148/5148            0s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-images-idx3-ubyte.gz
4422102/4422102      0s
0us/step
```

```
[3]: # Normalize pixel values to be between 0 and 1
      x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
[4]: # Step 3: Understand the structure of the dataset
      print(f"Training data shape: {x_train.shape}")
      print(f"Test data shape: {x_test.shape}")
      print(f"Number of classes: {len(set(y_train))}")
```

```
Training data shape: (60000, 28, 28)
Test data shape: (10000, 28, 28)
Number of classes: 10
```

```
[6]: # Step 4: Visualize the data
import matplotlib.pyplot as plt
class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal', 'Ankle boot', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

plt.figure(figsize=(10, 10))
for i in range(25):
    plt.subplot(5, 5, i + 1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i], cmap=plt.cm.binary)
    plt.xlabel(class_names[y_train[i]])
plt.show()
```



```
[7]: # Step 5: Form the Input, Hidden, and Output Layers
def create_model(activation_function):
    model = Sequential([
        Flatten(input_shape=(28, 28)), # Fashion MNIST input shape
        Dense(128, activation=activation_function),
        Dense(64, activation=activation_function),
        Dense(10, activation='softmax') # Output layer for classification
    ])
    return model
```

```
[8]: # Step 6: Compile the Model
def compile_model(model):
    model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
[9]: # Step 7: Fit the Model
def fit_model(model, x_train, y_train, epochs=10):
    history = model.fit(x_train, y_train, epochs=epochs, validation_split=0.2,
        ↪ verbose=2)
    return history
```

```
[14]: # Step 8: Find Accuracy of the Model
activation_functions = ['linear', 'sigmoid', 'tanh', 'relu', 'softmax']
accuracies = [] # Initialize a list to store accuracies

for activation_function in activation_functions:
    print(f"\nTraining with {activation_function} activation function")
    model = create_model(activation_function)
    compile_model(model)
    history = fit_model(model, x_train, y_train)
    test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
    print(f"Test accuracy with {activation_function}: {test_acc}")
    accuracies.append(test_acc) # Append the accuracy to the list
```

Training with linear activation function

```
/usr/local/lib/python3.10/dist-
packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
    super().__init__(**kwargs)
```

Epoch 1/10

1500/1500 - 7s - 5ms/step - accuracy: 0.8066 - loss: 0.5555 - val\_accuracy:



0.8334 - val\_loss: 0.4724  
Epoch 2/10  
1500/1500 - 7s - 5ms/step - accuracy: 0.8337 - loss: 0.4775 - val\_accuracy:  
0.8306 - val\_loss: 0.5015  
Epoch 3/10  
1500/1500 - 10s - 7ms/step - accuracy: 0.8420 - loss: 0.4552 - val\_accuracy:  
0.8423 - val\_loss: 0.4731  
Epoch 4/10  
1500/1500 - 6s - 4ms/step - accuracy: 0.8438 - loss: 0.4436 - val\_accuracy:  
0.8478 - val\_loss: 0.4363  
Epoch 5/10  
1500/1500 - 7s - 4ms/step - accuracy: 0.8473 - loss: 0.4370 - val\_accuracy:  
0.8465 - val\_loss: 0.4427  
Epoch 6/10  
1500/1500 - 7s - 5ms/step - accuracy: 0.8495 - loss: 0.4283 - val\_accuracy:  
0.8320 - val\_loss: 0.4703  
Epoch 7/10  
1500/1500 - 5s - 3ms/step - accuracy: 0.8529 - loss: 0.4232 - val\_accuracy:  
0.8533 - val\_loss: 0.4282  
Epoch 8/10  
1500/1500 - 10s - 7ms/step - accuracy: 0.8527 - loss: 0.4183 - val\_accuracy:  
0.8497 - val\_loss: 0.4328  
Epoch 9/10  
1500/1500 - 7s - 5ms/step - accuracy: 0.8557 - loss: 0.4145 - val\_accuracy:  
0.8488 - val\_loss: 0.4350  
Epoch 10/10  
1500/1500 - 6s - 4ms/step - accuracy: 0.8555 - loss: 0.4102 - val\_accuracy:  
0.8342 - val\_loss: 0.4688  
313/313 - 1s - 2ms/step - accuracy: 0.8255 - loss: 0.4976  
Test accuracy with linear: 0.8255000114440918

Training with sigmoid activation function

Epoch 1/10  
1500/1500 - 7s - 4ms/step - accuracy: 0.7811 - loss: 0.6618 - val\_accuracy:  
0.8402 - val\_loss: 0.4366  
Epoch 2/10  
1500/1500 - 7s - 5ms/step - accuracy: 0.8549 - loss: 0.4037 - val\_accuracy:  
0.8618 - val\_loss: 0.3849  
Epoch 3/10  
1500/1500 - 10s - 7ms/step - accuracy: 0.8705 - loss: 0.3624 - val\_accuracy:  
0.8633 - val\_loss: 0.3835  
Epoch 4/10  
1500/1500 - 8s - 5ms/step - accuracy: 0.8781 - loss: 0.3362 - val\_accuracy:  
0.8780 - val\_loss: 0.3428  
Epoch 5/10  
1500/1500 - 7s - 5ms/step - accuracy: 0.8844 - loss: 0.3176 - val\_accuracy:  
0.8789 - val\_loss: 0.3328  
Epoch 6/10

1500/1500 - 9s - 6ms/step - accuracy: 0.8898 - loss: 0.3014 - val\_accuracy:  
0.8842 - val\_loss: 0.3204  
Epoch 7/10  
1500/1500 - 10s - 7ms/step - accuracy: 0.8940 - loss: 0.2881 - val\_accuracy:  
0.8807 - val\_loss: 0.3273  
Epoch 8/10  
1500/1500 - 7s - 5ms/step - accuracy: 0.8991 - loss: 0.2759 - val\_accuracy:  
0.8857 - val\_loss: 0.3121  
Epoch 9/10  
1500/1500 - 10s - 7ms/step - accuracy: 0.9017 - loss: 0.2655 - val\_accuracy:  
0.8862 - val\_loss: 0.3151  
Epoch 10/10  
1500/1500 - 8s - 6ms/step - accuracy: 0.9070 - loss: 0.2528 - val\_accuracy:  
0.8873 - val\_loss: 0.3140  
313/313 - 1s - 2ms/step - accuracy: 0.8775 - loss: 0.3472  
Test accuracy with sigmoid: 0.8774999976158142

Training with tanh activation function

Epoch 1/10  
1500/1500 - 8s - 5ms/step - accuracy: 0.8252 - loss: 0.4871 - val\_accuracy:  
0.8521 - val\_loss: 0.4142  
Epoch 2/10  
1500/1500 - 6s - 4ms/step - accuracy: 0.8624 - loss: 0.3748 - val\_accuracy:  
0.8733 - val\_loss: 0.3503  
Epoch 3/10  
1500/1500 - 6s - 4ms/step - accuracy: 0.8758 - loss: 0.3362 - val\_accuracy:  
0.8684 - val\_loss: 0.3591  
Epoch 4/10  
1500/1500 - 6s - 4ms/step - accuracy: 0.8845 - loss: 0.3135 - val\_accuracy:  
0.8804 - val\_loss: 0.3297  
Epoch 5/10  
1500/1500 - 8s - 5ms/step - accuracy: 0.8886 - loss: 0.2978 - val\_accuracy:  
0.8791 - val\_loss: 0.3347  
Epoch 6/10  
1500/1500 - 10s - 7ms/step - accuracy: 0.8947 - loss: 0.2827 - val\_accuracy:  
0.8792 - val\_loss: 0.3312  
Epoch 7/10  
1500/1500 - 10s - 7ms/step - accuracy: 0.8986 - loss: 0.2739 - val\_accuracy:  
0.8747 - val\_loss: 0.3404  
Epoch 8/10  
1500/1500 - 6s - 4ms/step - accuracy: 0.9012 - loss: 0.2636 - val\_accuracy:  
0.8819 - val\_loss: 0.3286  
Epoch 9/10  
1500/1500 - 7s - 5ms/step - accuracy: 0.9053 - loss: 0.2528 - val\_accuracy:  
0.8856 - val\_loss: 0.3105  
Epoch 10/10  
1500/1500 - 9s - 6ms/step - accuracy: 0.9093 - loss: 0.2419 - val\_accuracy:  
0.8823 - val\_loss: 0.3244

313/313 - 1s - 2ms/step - accuracy: 0.8749 - loss: 0.3566  
Test accuracy with tanh: 0.8748999834060669

Training with relu activation function

Epoch 1/10

1500/1500 - 8s - 5ms/step - accuracy: 0.8154 - loss: 0.5189 - val\_accuracy:  
0.8486 - val\_loss: 0.4098

Epoch 2/10

1500/1500 - 11s - 7ms/step - accuracy: 0.8605 - loss: 0.3792 - val\_accuracy:  
0.8673 - val\_loss: 0.3720

Epoch 3/10

1500/1500 - 5s - 4ms/step - accuracy: 0.8749 - loss: 0.3420 - val\_accuracy:  
0.8695 - val\_loss: 0.3629

Epoch 4/10

1500/1500 - 11s - 7ms/step - accuracy: 0.8801 - loss: 0.3231 - val\_accuracy:  
0.8692 - val\_loss: 0.3642

Epoch 5/10

1500/1500 - 10s - 6ms/step - accuracy: 0.8886 - loss: 0.3012 - val\_accuracy:  
0.8743 - val\_loss: 0.3574

Epoch 6/10

1500/1500 - 7s - 4ms/step - accuracy: 0.8924 - loss: 0.2879 - val\_accuracy:  
0.8808 - val\_loss: 0.3278

Epoch 7/10

1500/1500 - 6s - 4ms/step - accuracy: 0.8994 - loss: 0.2726 - val\_accuracy:  
0.8695 - val\_loss: 0.3628

Epoch 8/10

1500/1500 - 7s - 5ms/step - accuracy: 0.9016 - loss: 0.2608 - val\_accuracy:  
0.8824 - val\_loss: 0.3270

Epoch 9/10

1500/1500 - 9s - 6ms/step - accuracy: 0.9055 - loss: 0.2515 - val\_accuracy:  
0.8808 - val\_loss: 0.3466

Epoch 10/10

1500/1500 - 7s - 5ms/step - accuracy: 0.9093 - loss: 0.2439 - val\_accuracy:  
0.8814 - val\_loss: 0.3421

313/313 - 1s - 2ms/step - accuracy: 0.8714 - loss: 0.3664

Test accuracy with relu: 0.871399998664856

Training with softmax activation function

Epoch 1/10

1500/1500 - 9s - 6ms/step - accuracy: 0.4897 - loss: 1.8767 - val\_accuracy:  
0.5660 - val\_loss: 1.2919

Epoch 2/10

1500/1500 - 9s - 6ms/step - accuracy: 0.5723 - loss: 1.0781 - val\_accuracy:  
0.5757 - val\_loss: 0.9840

Epoch 3/10

1500/1500 - 7s - 5ms/step - accuracy: 0.5817 - loss: 0.9277 - val\_accuracy:  
0.5893 - val\_loss: 0.9219

Epoch 4/10

```
1500/1500 - 11s - 7ms/step - accuracy: 0.5933 - loss: 0.8864 - val_accuracy:
0.6028 - val_loss: 0.8960
Epoch 5/10
1500/1500 - 5s - 4ms/step - accuracy: 0.6095 - loss: 0.8645 - val_accuracy:
0.6289 - val_loss: 0.8703
Epoch 6/10
1500/1500 - 11s - 7ms/step - accuracy: 0.7003 - loss: 0.7840 - val_accuracy:
0.7238 - val_loss: 0.7079
Epoch 7/10
1500/1500 - 7s - 5ms/step - accuracy: 0.7421 - loss: 0.6391 - val_accuracy:
0.7457 - val_loss: 0.6274
Epoch 8/10
1500/1500 - 6s - 4ms/step - accuracy: 0.7533 - loss: 0.5874 - val_accuracy:
0.7587 - val_loss: 0.5843
Epoch 9/10
1500/1500 - 8s - 5ms/step - accuracy: 0.7766 - loss: 0.5479 - val_accuracy:
0.7747 - val_loss: 0.5636
Epoch 10/10
1500/1500 - 8s - 6ms/step - accuracy: 0.7902 - loss: 0.5157 - val_accuracy:
0.7793 - val_loss: 0.5474
313/313 - 1s - 3ms/step - accuracy: 0.7773 - loss: 0.5580
Test accuracy with softmax: 0.7773000001907349
```

```
[15]: # Plot the graph comparing activation functions' accuracy
plt.figure(figsize=(10, 6))
plt.bar(activation_functions, accuracies, color='skyblue')
plt.xlabel('Activation Functions')
plt.ylabel('Accuracy (%)')
plt.title('Comparison of Activation Functions Accuracy')
plt.ylim([0, 1]) # Set y-axis range from 0 to 1 (accuracy percentage)
plt.show()
```

