

Experiment No. 6																				
BE (AI&DS)		ROLL NO : 9742																		
Date of Implementation: 06/09/2024																				
Aim: Implement auto encoder for image compression																				
Programming Language Used : Python																				
Upon completion of this experiment, students will be able to LO3: Build and train deep learning models for given problem																				
<table border="1"> <tr> <td>Indicator</td> <td></td> <td></td> </tr> <tr> <td>Timeline</td> <td></td> <td></td> </tr> <tr> <td>Maintains submission deadline (1)</td> <td>On time (1)</td> <td>Otherwise (0)</td> </tr> <tr> <td>Completion and Organization (2)</td> <td>Completed in LAB (2)</td> <td>Otherwise(1)</td> </tr> <tr> <td>Analysis of output and conclusion(2)</td> <td>Properly done (2)</td> <td>Otherwise (0)</td> </tr> <tr> <td>Viva (10)</td> <td></td> <td></td> </tr> </table>			Indicator			Timeline			Maintains submission deadline (1)	On time (1)	Otherwise (0)	Completion and Organization (2)	Completed in LAB (2)	Otherwise(1)	Analysis of output and conclusion(2)	Properly done (2)	Otherwise (0)	Viva (10)		
Indicator																				
Timeline																				
Maintains submission deadline (1)	On time (1)	Otherwise (0)																		
Completion and Organization (2)	Completed in LAB (2)	Otherwise(1)																		
Analysis of output and conclusion(2)	Properly done (2)	Otherwise (0)																		
Viva (10)																				
Assessment Marks : <table border="1"> <tr> <td>Timeline(1)</td> <td></td> </tr> <tr> <td>Completion and Organization (2)</td> <td></td> </tr> <tr> <td>Analysis of output and conclusion(2)</td> <td></td> </tr> <tr> <td>Viva (10)</td> <td></td> </tr> <tr> <td>Total (15)</td> <td></td> </tr> </table>			Timeline(1)		Completion and Organization (2)		Analysis of output and conclusion(2)		Viva (10)		Total (15)									
Timeline(1)																				
Completion and Organization (2)																				
Analysis of output and conclusion(2)																				
Viva (10)																				
Total (15)																				

EXPERIMENT	6
Aim	To Implement Auto encoder for Image Compression
Tools	PYTHON
Theory	<p>Traditional feedforward neural networks can be great at performing tasks such as classification and regression, but what if we would like to implement solutions such as signal denoising or anomaly detection? One way to do this is by using Autoencoders. Autoencoders are a specialized class of algorithms that can learn efficient representations of input data with no need for labels. It is a class of artificial neural networks designed for unsupervised learning. Learning to compress and effectively represent input data without specific labels is the essential principle of an automatic decoder. This is accomplished using a two-fold structure that consists of an encoder and a decoder. The encoder transforms the input data into a reduced-dimensional representation, which is often referred to as “latent space” or “encoding”. From that representation, a decoder rebuilds the initial input. For the network to gain meaningful patterns in data, a process of encoding and decoding facilitates the definition of essential features.</p> <p>Architecture of Autoencoder in Deep Learning</p> <p>The general architecture of an autoencoder includes an encoder, decoder, and bottleneck layer.</p> <div data-bbox="449 1068 1029 1617" data-label="Diagram"> <p>The diagram illustrates the architecture of an autoencoder. It consists of three layers: an Input layer, a Hidden layer, and an Output layer. The Input layer has 6 nodes labeled $X_1, X_2, X_3, X_4, X_5, X_6$. The Hidden layer has 3 nodes labeled a_1, a_2, a_3, with the text "bottleneck" written above them. The Output layer has 6 nodes labeled $\hat{X}_1, \hat{X}_2, \hat{X}_3, \hat{X}_4, \hat{X}_5, \hat{X}_6$. All nodes in the Input layer are connected to all nodes in the Hidden layer, and all nodes in the Hidden layer are connected to all nodes in the Output layer, forming a fully connected feedforward network.</p> </div> <p>1. Encoder</p> <ul style="list-style-type: none"> • Input layer take raw input data • The hidden layers progressively reduce the dimensionality of the input, capturing important features and patterns. These layer compose the encoder.

	<ul style="list-style-type: none"> • The bottleneck layer (latent space) is the final hidden layer, where the dimensionality is significantly reduced. This layer represents the compressed encoding of the input data. <ol style="list-style-type: none"> 2. Decoder <ul style="list-style-type: none"> • The bottleneck layer takes the encoded representation and expands it back to the dimensionality of the original input. • The hidden layers progressively increase the dimensionality and aim to reconstruct the original input. • The output layer produces the reconstructed output, which ideally should be as close as possible to the input data. 3. The loss function used during training is typically a reconstruction loss, measuring the difference between the input and the reconstructed output. Common choices include mean squared error (MSE) for continuous data or binary cross-entropy for binary data. 4. During training, the autoencoder learns to minimize the reconstruction loss, forcing the network to capture the most important features of the input data in the bottleneck layer.
Implementation	<ol style="list-style-type: none"> 1. Import necessary libraries 2. Load mnist digit data set 3. Define a basic autoencoder 4. Compile and fit Autoencoder 5. Visualize the original and reconstructed data
Conclusion	<p>In this experiment, we implemented an Autoencoder for image compression using the MNIST dataset. The model learned to compress images by reducing dimensionality to a latent space and successfully reconstructed the original images. The training process showed decreasing loss values, indicating effective learning. Visualization confirmed the model's ability to reproduce the input data. Overall, this experiment highlighted the potential of Autoencoders as powerful tools for unsupervised learning and applications such as image denoising and anomaly detection.</p>

Implementation:

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist
```

```
[2]: # Load MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()

# Normalize the data (Pixel values between 0 and 1)
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Reshape to fit into the model (flattened into 784-dimensional vectors)
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 0s
0us/step

```
[3]: # Define the size of the encoded representation (latent space size)
encoding_dim = 32

# Input placeholder
input_img = Input(shape=(784,))

# Encoder: Encodes input into a smaller representation
encoded = Dense(encoding_dim, activation='relu')(input_img)

# Decoder: Decodes the smaller representation back to the original input size
decoded = Dense(784, activation='sigmoid')(encoded)

# Model that maps input to its reconstruction
autoencoder = Model(input_img, decoded)

# Separate Encoder Model
```

```

encoder = Model(input_img, encoded)

# Create a decoder model (decoder layer)
encoded_input = Input(shape=(encoding_dim,))
decoder_layer = autoencoder.layers[-1] # Last layer of the autoencoder
decoder = Model(encoded_input, decoder_layer(encoded_input))

```

```

[6]: # Compile the autoencoder
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Train the autoencoder and store the training history
history = autoencoder.fit(x_train, x_train,
                          epochs=50,
                          batch_size=256,
                          shuffle=True,
                          validation_data=(x_test, x_test))

```

```

Epoch 1/50
235/235          5s 17ms/step -
loss: 0.0926 - val_loss: 0.0915
Epoch 2/50
235/235          3s 11ms/step -
loss: 0.0925 - val_loss: 0.0915
Epoch 3/50
235/235          5s 13ms/step -
loss: 0.0925 - val_loss: 0.0915
Epoch 4/50
235/235          3s 13ms/step -
loss: 0.0925 - val_loss: 0.0915
Epoch 5/50
235/235          4s 10ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 6/50
235/235          2s 9ms/step -
loss: 0.0925 - val_loss: 0.0915
Epoch 7/50
235/235          3s 9ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 8/50
235/235          3s 11ms/step -
loss: 0.0923 - val_loss: 0.0915
Epoch 9/50
235/235          5s 9ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 10/50
235/235          3s 9ms/step -
loss: 0.0925 - val_loss: 0.0914

```

Epoch 11/50
235/235 2s 9ms/step -
loss: 0.0926 - val_loss: 0.0914
Epoch 12/50
235/235 3s 11ms/step -
loss: 0.0923 - val_loss: 0.0914
Epoch 13/50
235/235 5s 9ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 14/50
235/235 2s 9ms/step -
loss: 0.0923 - val_loss: 0.0914
Epoch 15/50
235/235 3s 13ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 16/50
235/235 4s 16ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 17/50
235/235 4s 15ms/step -
loss: 0.0923 - val_loss: 0.0914
Epoch 18/50
235/235 2s 9ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 19/50
235/235 2s 9ms/step -
loss: 0.0922 - val_loss: 0.0914
Epoch 20/50
235/235 3s 9ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 21/50
235/235 4s 14ms/step -
loss: 0.0924 - val_loss: 0.0913
Epoch 22/50
235/235 4s 9ms/step -
loss: 0.0922 - val_loss: 0.0914
Epoch 23/50
235/235 2s 9ms/step -
loss: 0.0925 - val_loss: 0.0915
Epoch 24/50
235/235 2s 9ms/step -
loss: 0.0925 - val_loss: 0.0914
Epoch 25/50
235/235 3s 11ms/step -
loss: 0.0922 - val_loss: 0.0914
Epoch 26/50
235/235 5s 10ms/step -
loss: 0.0924 - val_loss: 0.0913

Epoch 27/50
235/235 2s 10ms/step -
loss: 0.0924 - val_loss: 0.0914
Epoch 28/50
235/235 3s 10ms/step -
loss: 0.0923 - val_loss: 0.0913
Epoch 29/50
235/235 3s 11ms/step -
loss: 0.0923 - val_loss: 0.0913
Epoch 30/50
235/235 5s 10ms/step -
loss: 0.0925 - val_loss: 0.0913
Epoch 31/50
235/235 2s 9ms/step -
loss: 0.0923 - val_loss: 0.0913
Epoch 32/50
235/235 2s 9ms/step -
loss: 0.0922 - val_loss: 0.0913
Epoch 33/50
235/235 3s 11ms/step -
loss: 0.0923 - val_loss: 0.0913
Epoch 34/50
235/235 3s 13ms/step -
loss: 0.0924 - val_loss: 0.0913
Epoch 35/50
235/235 4s 10ms/step -
loss: 0.0925 - val_loss: 0.0913
Epoch 36/50
235/235 2s 9ms/step -
loss: 0.0922 - val_loss: 0.0913
Epoch 37/50
235/235 3s 10ms/step -
loss: 0.0924 - val_loss: 0.0913
Epoch 38/50
235/235 5s 21ms/step -
loss: 0.0920 - val_loss: 0.0913
Epoch 39/50
235/235 2s 9ms/step -
loss: 0.0923 - val_loss: 0.0912
Epoch 40/50
235/235 2s 9ms/step -
loss: 0.0922 - val_loss: 0.0913
Epoch 41/50
235/235 3s 10ms/step -
loss: 0.0923 - val_loss: 0.0913
Epoch 42/50
235/235 2s 9ms/step -
loss: 0.0924 - val_loss: 0.0912

```
Epoch 43/50
235/235          4s 15ms/step -
loss: 0.0922 - val_loss: 0.0914
Epoch 44/50
235/235          4s 10ms/step -
loss: 0.0923 - val_loss: 0.0913
Epoch 45/50
235/235          2s 9ms/step -
loss: 0.0922 - val_loss: 0.0913
Epoch 46/50
235/235          3s 9ms/step -
loss: 0.0922 - val_loss: 0.0913
Epoch 47/50
235/235          4s 13ms/step -
loss: 0.0924 - val_loss: 0.0913
Epoch 48/50
235/235          3s 12ms/step -
loss: 0.0924 - val_loss: 0.0913
Epoch 49/50
235/235          5s 9ms/step -
loss: 0.0925 - val_loss: 0.0912
Epoch 50/50
235/235          2s 9ms/step -
loss: 0.0923 - val_loss: 0.0912
```

```
[7]: # Encode and decode some digits from the test set
    encoded_imgs = encoder.predict(x_test)
    decoded_imgs = decoder.predict(encoded_imgs)


    # Visualize the original and reconstructed images
    n = 10 # Number of digits to display
    plt.figure(figsize=(20, 4))

    for i in range(n):
        # Display original images
        ax = plt.subplot(2, n, i + 1)
        plt.imshow(x_test[i].reshape(28, 28), cmap='gray')
        plt.title("Original")
        plt.axis('off')


        # Display reconstructed images
        ax = plt.subplot(2, n, i + 1 + n)
        plt.imshow(decoded_imgs[i].reshape(28, 28), cmap='gray')
        plt.title("Reconstructed")
        plt.axis('off')

    plt.show()
```


313/313 1s 2ms/step
313/313 0s 1ms/step



```
[8]: # Plot training history
plt.figure(figsize=(10, 4))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Autoencoder Training History')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
[10]: # Calculate and print compression ratio
original_size = x_test.nbytes
encoded_imgs = encoder.predict(x_test) # Get the encoded images
compressed_size = encoded_imgs.nbytes
compression_ratio = original_size / compressed_size
print(f"Compression Ratio: {compression_ratio:.2f}")
```

313/313 0s 1ms/step

Compression Ratio: 24.50